

# Repeatable splitting

In this notebook, we will explore the impact of different ways of creating machine learning datasets.

Repeatability is important in machine learning. If you do the same thing now and 5 minutes from now and get different answers, then it makes experimentation difficult. In other words, you will find it difficult to gauge whether a change you made has resulted in an improvement or not.

```
In [1]: !sudo chown -R jupyter:jupyter /home/jupyter/training-data-analyst
```

```
In [2]: !pip install --user google-cloud-bigquery==1.25.0
```

Collecting google-cloud-bigquery==1.25.0

Downloading google\_cloud\_bigquery-1.25.0-py2.py3-none-any.whl (169 kB)

|██| 169 kB 5.2 MB/s eta 0:00:01

Requirement already satisfied: six<2.0.0dev,>=1.13.0 in /opt/conda/lib/python3.7/site-packages (from google-cloud-bigquery==1.25.0) (1.16.0)

Requirement already satisfied: google-auth<2.0dev,>=1.9.0 in /opt/conda/lib/python3.7/site-packages (from google-cloud-bigquery==1.25.0) (1.34.0)

Collecting google-resumable-media<0.6dev,>=0.5.0

Downloading google\_resumable\_media-0.5.1-py2.py3-none-any.whl (38 kB)

Requirement already satisfied: google-cloud-core<2.0dev,>=1.1.0 in /opt/conda/lib/python3.7/site-packages (from google-cloud-bigquery==1.25.0) (1.7.2)

Requirement already satisfied: google-api-core<2.0dev,>=1.15.0 in /opt/conda/lib/python3.7/site-packages (from google-cloud-bigquery==1.25.0) (1.31.1)

Requirement already satisfied: protobuf<=3.6.0 in /opt/conda/lib/python3.7/site-packages (from google-cloud-bigquery==1.25.0) (3.16.0)

Requirement already satisfied: packaging<=14.3 in /opt/conda/lib/python3.7/site-packages (from google-api-core<2.0dev,>=1.15.0->google-cloud-bigquery==1.25.0) (21.0)

Requirement already satisfied: googleapis-common-protos<2.0dev,>=1.6.0 in /opt/conda/lib/python3.7/site-packages (from google-api-core<2.0dev,>=1.15.0->google-cloud-bigquery==1.25.0) (1.53.0)

Requirement already satisfied: requests<3.0.0dev,>=2.18.0 in /opt/conda/lib/python3.7/site-packages (from google-api-core<2.0dev,>=1.15.0->google-cloud-bigquery==1.25.0) (2.25.1)

Requirement already satisfied: pytz in /opt/conda/lib/python3.7/site-packages (from google-api-core<2.0dev,>=1.15.0->google-cloud-bigquery==1.25.0) (2021.1)

Requirement already satisfied: setuptools<=40.3.0 in /opt/conda/lib/python3.7/site-packages (from google-api-core<2.0dev,>=1.15.0->google-cloud-bigquery==1.25.0) (49.6.0.post20210108)

Requirement already satisfied: pyasn1-modules<=0.2.1 in /opt/conda/lib/python3.7/site-packages (from google-auth<2.0dev,>=1.9.0->google-cloud-bigquery==1.25.0) (0.2.7)

Requirement already satisfied: cachetools<5.0,>=2.0.0 in /opt/conda/lib/python3.7/site-packages (from google-auth<2.0dev,>=1.9.0->google-cloud-bigquery==1.25.0) (4.2.2)

Requirement already satisfied: rsa<5,>=3.1.4 in /opt/conda/lib/python3.7/site-packages (from google-auth<2.0dev,>=1.9.0->google-cloud-bigquery==1.25.0) (4.7.2)

Requirement already satisfied: pyparsing<=2.0.2 in /opt/conda/lib/python3.7/site-packages (from packaging<=14.3->google-api-core<2.0dev,>=1.15.0->google-cloud-bigquery==1.25.0) (2.4.7)

Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /opt/conda/lib/python3.7/site-packages (from pyasn1-modules>=0.2.1->google-auth<2.0dev,>=1.9.0->google-cloud-bigquery==1.25.0) (0.4.8)

Requirement already satisfied: certifi>=2017.4.17 in /opt/conda/lib/python3.7/site-packages (from requests<3.0.0dev,>=2.18.0->google-api-core<2.0dev,>=1.15.0->google-cloud-bigquery==1.25.0) (2021.5.30)

Requirement already satisfied: idna<3,>=2.5 in /opt/conda/lib/python3.7/site-packages (from requests<3.0.0dev,>=2.18.0->google-api-core<2.0dev,>=1.15.0->google-cloud-bigquery==1.25.0) (2.10)

Requirement already satisfied: chardet<5,>=3.0.2 in /opt/conda/lib/python3.7/site-packages (from requests<3.0.0dev,>=2.18.0->google-api-core<2.0dev,>=1.15.0->google-cloud-bigquery==1.25.0) (4.0.0)

Requirement already satisfied: urllib3<1.27,>=1.21.1 in /opt/conda/lib/python3.7/site-packages (from requests<3.0.0dev,>=2.18.0->google-api-core<2.0dev,>=1.15.0->google-cloud-bigquery==1.25.0) (1.26.6)

Installing collected packages: google-resumable-media, google-cloud-bigquery

ERROR: pip's dependency resolver does not currently take into account all the packages that are installed. This behaviour is the source of the following dependency conflicts.

tfx-bsl 1.2.0 requires absl-py<0.13,>=0.9, but you have absl-py 0.13.0 which is incompatible.

tfx-bsl 1.2.0 requires google-api-python-client<2,>=1.7.11, but you have google-api-python-client 2.15.0 which is incompatible.

tfx-bsl 1.2.0 requires google-cloud-bigquery<2.21,>=1.28.0, but you have google-cloud-bigquery 1.25.0 which is incompatible.

tfx-bsl 1.2.0 requires pyarrow<3,>=1, but you have pyarrow 5.0.0 which is incompatible.

tensorflow-transform 1.2.0 requires absl-py<0.13,>=0.9, but you have absl-py 0.13.0 which is incompatible.

tensorflow-transform 1.2.0 requires google-cloud-bigquery<2.21,>=1.28.0, but you have google-cloud-bigquery 1.25.0 which is incompatible.

tensorflow-transform 1.2.0 requires pyarrow<3,>=1, but you have pyarrow 5.0.0 which is incompatible.

google-cloud-storage 1.41.1 requires google-resumable-media<3.0dev,>=1.3.0; python\_version >= "3.6", but you have google-resumable-media 0.5.1 which is incompatible.

Successfully installed google-cloud-bigquery-1.25.0 google-resumable-media-0.5.1

**Restart** the kernel before proceeding further (On the Notebook menu - Kernel - Restart Kernel).

In [3]: `from google.cloud import bigquery`

## Create a simple machine learning model

The dataset that we will use is [a BigQuery public dataset](#) of airline arrival data. Click on the link, and look at the column names. Switch to the Details tab to verify that the number of records is 70 million, and then switch to the Preview tab to look at a few rows.

We want to predict the arrival delay of an airline based on the departure delay. The model that we will use is a zero-bias linear model:  $\text{delay}_{\text{arrival}} = \alpha * \text{delay}_{\text{departure}}$

To train the model is to estimate a good value for  $\alpha$ .

One approach to estimate  $\alpha$  is to use this formula:  $\alpha = \frac{\sum \text{delay}_{\text{departure}} \text{delay}_{\text{arrival}}}{\sum \text{delay}_{\text{departure}}^2}$  Because we'd like to

capture the idea that this relationship is different for flights from New York to Los Angeles vs. flights from Austin to Indianapolis (shorter flight, less busy airports), we'd compute a different  $\alpha$  for each airport-pair. For simplicity, we'll do this model only for flights between Denver and Los Angeles.

## Naive random split (not repeatable)

In [4]:

```
compute_alpha = """
#standardSQL
SELECT
  SAFE_DIVIDE(
    SUM(arrival_delay * departure_delay),
    SUM(departure_delay * departure_delay)) AS alpha
FROM
  (
    SELECT
      RAND() AS splitfield,
      arrival_delay,
      departure_delay
    FROM
      `bigquery-samples.airline_ontime_data.flights`
    WHERE
      departure_airport = 'DEN'
      AND arrival_airport = 'LAX'
  )
WHERE
  splitfield < 0.8
"""
```

In [5]:

```
results = bigquery.Client().query(compute_alpha).to_dataframe()
alpha = results['alpha'][0]
print(alpha)
```

0.9746759869400119

What is wrong with calculating RMSE on the training and test data as follows?

In [6]:

```

compute_rmse = """
#standardSQL
SELECT
  dataset,
  SQRT(
    AVG(
      (arrival_delay - ALPHA * departure_delay) *
      (arrival_delay - ALPHA * departure_delay)
    )
  ) AS rmse,
  COUNT(arrival_delay) AS num_flights
FROM (
  SELECT
    IF (RAND() < 0.8, 'train', 'eval') AS dataset,
    arrival_delay,
    departure_delay
  FROM
    `bigquery-samples.airline_ontime_data.flights`
  WHERE
    departure_airport = 'DEN'
    AND arrival_airport = 'LAX' )
GROUP BY
  dataset
"""
bigquery.Client().query(compute_rmse.replace('ALPHA', str(alpha))).to_dataframe()

```

Out[6]:

	dataset	rmse	num_flights
0	eval	13.060312	16042
1	train	13.089737	63647

Hint:

- Are you really getting the same training data in the compute\_rmse query as in the compute\_alpha query?
- Do you get the same answers each time you rerun the compute\_alpha and compute\_rmse blocks?

## How do we correctly train and evaluate?

Here's the right way to compute the RMSE using the actual training and held-out (evaluation) data. Note how much harder this feels.

Although the calculations are now correct, the experiment is still not repeatable.

Try running it several times; do you get the same answer?

```
In [7]: train_and_eval_rand = """
#standardSQL
WITH
  alldata AS (
    SELECT
      IF (RAND() < 0.8, 'train', 'eval') AS dataset,
      arrival_delay,
      departure_delay
    FROM
      `bigquery-samples.airline_ontime_data.flights`
    WHERE
      departure_airport = 'DEN'
      AND arrival_airport = 'LAX' ),
  training AS (
    SELECT
      SAFE_DIVIDE(
        SUM(arrival_delay * departure_delay),
        SUM(departure_delay * departure_delay)) AS alpha
    FROM
      alldata
    WHERE
      dataset = 'train' )

SELECT
  MAX(alpha) AS alpha,
  dataset,
  SQRT(
    AVG(
      (arrival_delay - alpha * departure_delay) *
      (arrival_delay - alpha * departure_delay)
    )
  ) AS rmse,
  COUNT(arrival_delay) AS num_flights
FROM
  alldata,
  training
GROUP BY
  dataset
"""
```

```
In [8]: bigquery.Client().query(train_and_eval_rand).to_dataframe()
```

```
Out[8]:
```

	alpha	dataset	rmse	num_flights
0	0.975873	train	13.023342	63689
1	0.975873	eval	13.321639	16000

## Using HASH of date to split the data

Let's split by date and train.

In [9]:

```

compute_alpha = """
#standardSQL
SELECT
    SAFE_DIVIDE(
        SUM(arrival_delay * departure_delay),
        SUM(departure_delay * departure_delay)) AS alpha
FROM
    `bigquery-samples.airline_ontime_data.flights`
WHERE
    departure_airport = 'DEN'
    AND arrival_airport = 'LAX'
    AND ABS(MOD(FARM_FINGERPRINT(date), 10)) < 8
"""
results = bigquery.Client().query(compute_alpha).to_dataframe()
alpha = results['alpha'][0]
print(alpha)

```

0.9758039143620403

We can now use the alpha to compute RMSE. Because the alpha value is repeatable, we don't need to worry that the alpha in the compute\_rmse will be different from the alpha computed in the compute\_alpha.

In [10]:

```

compute_rmse = """
#standardSQL
SELECT
    IF(ABS(MOD(FARM_FINGERPRINT(date), 10)) < 8, 'train', 'eval') AS dataset,
    SQRT(
        AVG(
            (arrival_delay - ALPHA * departure_delay) *
            (arrival_delay - ALPHA * departure_delay)
        )
    ) AS rmse,
    COUNT(arrival_delay) AS num_flights
FROM
    `bigquery-samples.airline_ontime_data.flights`
WHERE
    departure_airport = 'DEN'
    AND arrival_airport = 'LAX'
GROUP BY
    dataset
"""
print(bigquery.Client().query(compute_rmse.replace('ALPHA', str(alpha))).to_dataframe()

```

	dataset	rmse	num_flights
0	eval	12.764685	15671
1	train	13.160712	64018

Note also that the RMSE on the evaluation dataset more from the RMSE on the training dataset when we do the split correctly. This should be expected; in the RAND() case, there was leakage between training and evaluation datasets, because there is high correlation between flights on the same day.

This is one of the biggest dangers with doing machine learning splits the wrong way -- **you will develop a false sense of confidence in how good your model is!**

Copyright 2018 Google Inc. Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the

License at <http://www.apache.org/licenses/LICENSE-2.0> Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License